

Graziano Montanaro

Data Engineer

Italy · EU citizen · Open to relocation/remote

✉ graziano.montanaro98@gmail.com

[in](#) [graziano-montanaro](#)

[montanarograziano](#)

[Google Scholar](#)

Data engineer building distributed data platforms, production APIs, and reliable lakehouse pipelines. Open-source contributor and speaker with a research background in machine learning.

Experience

09/2024– Present **Data Engineer, Intella**, Milan (Remote)

- Set company-wide data engineering standards and built shared Python libraries used across five microservices by data, ML, and full-stack engineers, covering configuration, structured logging, storage I/O, and data quality.
- Re-architected the storage layer: migrated terabyte-scale datasets to Delta Lake and rebuilt Prefect orchestration on Kubernetes, cutting query latency up to 90% and memory use via predicate and projection pushdown in Polars.
- Delivered the API and Prefect synchronization layer for LTTB-based telemetry downsampling on TimescaleDB, enabling interactive charting over multi-year ranges.
- Built a fault-tolerant microservice orchestrating hundreds of concurrent jobs within SLA, with retries and fault isolation.

02/2023– 09/2024 **Data Engineer, Tuidi**, Putignano

- Optimized ETL over terabyte-scale data by combining distributed PySpark on Databricks with single-node Polars for smaller workloads, cutting runtime by up to 70% on the heaviest pipelines.
- Introduced automated data quality checks and monitoring, improving reliability across downstream forecasting pipelines.
- Built a Dash dashboard for retail sales forecasting models, surfacing prediction quality, degradation, and target drift.

Skills

Programming & APIs Python, SQL, Java, FastAPI, Pydantic, Apache Arrow/Flight, gRPC

Data & Lakehouse PySpark, Databricks, Delta Lake, Iceberg, dbt, Airflow, Prefect, Pandas, Polars, DuckDB, DuckLake, Ibis, Great Expectations

ML & MLOps PyTorch, TensorFlow, scikit-learn, MLflow, DVC, SHAP, Grad-CAM; time series forecasting, computer vision

Streaming & Databases Kafka, RabbitMQ, PostgreSQL, TimescaleDB, MySQL, MongoDB, Redis

Cloud & DevOps AWS (S3, Glue), Docker, Kubernetes, GitLab CI/CD, GitHub Actions, Grafana, Git

Languages Italian (Native), English (Fluent)

Projects

Italy Dashboard, [GitHub](#)

Built a reproducible public-data platform unifying official Italian statistics from ISTAT, INPS, MUR, and weather providers.

- Implemented dbt, DuckDB, and Parquet pipelines with Python/Polars ingestion and browser-side DuckDB-WASM analytics, backed by CI-enforced provenance and frontend conformance checks.

Poor Man's Lakehouse, [GitHub](#)

Built a composable open-source lakehouse spanning local and AWS deployments.

- Unified four catalog backends, Iceberg and Delta Lake, and Spark, Polars, DuckDB, and Ibis access through a tested Python connector.

Education

09/2020– 07/2022 **Master's Degree in Computer Science and Artificial Intelligence**, University of Bari Aldo Moro, Bari, 110/110 cum laude (highest honours)

Multimodal 3D CNN thesis for Alzheimer's detection from MRI and amyloid PET, reaching 95% accuracy on OASIS-3.

09/2017– 07/2020 **Bachelor's Degree in Computer Science**, University of Bari Aldo Moro, Bari, 110/110 cum laude (highest honours)

Publications

2024 Automated detection of Alzheimer's disease: a multi-modal approach with 3D MRI and amyloid PET. *Scientific Reports*, 14, 5210.

2021 Detection of Dementia through 3D convolutional neural networks based on amyloid PET. *IEEE SSCL*.